

Predicting NANSE Store Traffic Through Generalized Linear Models (GLM)

Creating the confusion matrix

```
# RStudio Options
options(scipen = 1000) #Prevent display in scientific notation
# Run this reusable confusion matrix function (https://en.wikipedia.org/wiki/Confusion_matrix)
my_confusion_matrix <- function(cf_table) {
  true_positive <- cf_table[4]
  true_negative <- cf_table[1]
  false_positive <- cf_table[2]
  false_negative <- cf_table[3]
  accuracy <- (true_positive + true_negative) / (true_positive + true_negative + false_positive + false_negative)
  sensitivity_recall <- true_positive / (true_positive + false_negative)
  specificity_selectivity <- true_negative / (true_negative + false_positive)
  precision <- true_positive / (true_positive + false_positive)
  neg_pred_value <- true_negative / (true_negative + false_negative)
  print(cf_table)
  my_list <- list(sprintf("%1.0f = True Positive (TP), Hit", true_positive),
                  sprintf("%1.0f = True Negative (TN), Rejection", true_negative),
                  sprintf("%1.0f = False Positive (FP), Type 1 Error", false_positive),
                  sprintf("%1.0f = False Negative (FN), Type 2 Error", false_negative),
                  sprintf("%1.4f = Accuracy (TP+TN/(TP+TN+FP+FN))", accuracy),
                  sprintf("%1.4f = Sensitivity, Recall, Hit Rate, True Positive Rate (How many positives did the model correctly predict)", sensitivity_recall),
                  sprintf("%1.4f = Specificity, Selectivity, True Negative Rate (How many negatives did the model correctly predict)", specificity_selectivity),
                  sprintf("%1.4f = Precision, Positive Predictive Value (How good are the model's positive predictions)", precision),
                  sprintf("%1.4f = Negative Predictive Value (How good are the model's negative predictions)", neg_pred_value))
  return(my_list)
}
```

Installing and Loading Packages

```
#install.packages('tidyverse')
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr    1.5.0
## v ggplot2    3.4.4      v tibble     3.2.1
## v lubridate  1.9.2      v tidyr      1.3.0
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```

# Load data
df <- read_rds("mod6HE_logit.rds")

# Explore the data and discuss in PowerPoint
summary(df)

##      store      week      high_med_gp      high_med_rev
## 186      : 52   Min.      : 1.00   Min.      :0.0000   Min.      :0.0000
## 227      : 52   1st Qu.:13.00   1st Qu.:0.0000   1st Qu.:0.0000
## 233      : 52   Median :26.00   Median :0.0000   Median :0.0000
## 236      : 52   Mean    :26.47   Mean    :0.4997   Mean    :0.4997
## 272      : 52   3rd Qu.:39.00   3rd Qu.:1.0000   3rd Qu.:1.0000
## 291      : 52   Max.     :52.00   Max.     :1.0000   Max.     :1.0000
## (Other):9759
## high_med_units  high_med_gpm      size      region
## Min.      :0.0000   Min.      :0.0000   Min.      : 890.0   ONTARIO :3120
## 1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.: 916.0   WEST     :4776
## Median :0.0000   Median :0.0000   Median : 943.0   QUEBEC  :1863
## Mean      :0.4995   Mean      :0.4639   Mean      : 949.6   ATLANTIC: 312
## 3rd Qu.:1.0000   3rd Qu.:1.0000   3rd Qu.: 972.0
## Max.      :1.0000   Max.      :1.0000   Max.      :1163.0
##
## promo_units_per altbev_units_per confect_units_per salty_units_per
## Min.      :0.1053   Min.      :0.2030   Min.      :0.0000   Min.      :0.0000
## 1st Qu.:0.3046   1st Qu.:0.3575   1st Qu.:0.2349   1st Qu.:0.1520
## Median :0.3451   Median :0.4019   Median :0.2645   Median :0.1716
## Mean      :0.3506   Mean      :0.4055   Mean      :0.2706   Mean      :0.1737
## 3rd Qu.:0.3929   3rd Qu.:0.4491   3rd Qu.:0.3012   3rd Qu.:0.1928
## Max.      :0.5797   Max.      :0.7250   Max.      :0.4741   Max.      :0.3684
##
## velocityA_units_per velocityB_units_per velocityC_units_per
## Min.      :0.4429   Min.      :0.1410   Min.      :0.00000
## 1st Qu.:0.5744   1st Qu.:0.2032   1st Qu.:0.07078
## Median :0.6107   Median :0.2162   Median :0.07994
## Mean      :0.6086   Mean      :0.2169   Mean      :0.08009
## 3rd Qu.:0.6428   3rd Qu.:0.2300   3rd Qu.:0.08929
## Max.      :0.7500   Max.      :0.3135   Max.      :0.14865
##
## velocityD_units_per velocityNEW_units_per
## Min.      :0.00000   Min.      :0.000000
## 1st Qu.:0.05533   1st Qu.:0.000000
## Median :0.07045   Median :0.000000
## Mean      :0.07918   Mean      :0.004608
## 3rd Qu.:0.10265   3rd Qu.:0.007963
## Max.      :0.21637   Max.      :0.055623
##

```

Preparing the data for the logistic regression algorithm

```

# Not for the model (for use later)
ColumnsNotUsed <- df %>%
  select(store, week, high_med_rev, high_med_gpm, high_med_gp)

```

```
# For use in the model
logit1 <- df %>%
  select(high_med_units,size, region, promo_units_per,
         altbev_units_per, confect_units_per, salty_units_per,
         velocityA_units_per, velocityB_units_per, velocityC_units_per, velocityD_units_per, velocityNEW_units_per)
```

Partitioning the data into testing and training datasets

```
#install.packages('caret') (don't install twice)
library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

set.seed(42)
partition <- caret::createDataPartition(y=logit1$high_med_units, p=.75, list=FALSE) #gives matrix of rows
data_train <- logit1[partition, ] #keeps the rows indicated in `partition` and all columns from `logit1`
data_test <- logit1[-partition, ] #keeps the rows not indicated in `partition` and all columns from `logit1`
```

Training the multivariate model

These are the instructions part of machine learning

```
model_train <- glm(high_med_units ~ ., family=binomial, data=data_train)
summary(model_train)

##
## Call:
## glm(formula = high_med_units ~ ., family = binomial, data = data_train)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -33.7554575   3.4503718  -9.783 < 0.0000000000000002 ***
## size           0.0198825   0.0008586  23.156 < 0.0000000000000002 ***
## regionWEST     1.2372452   0.0747641  16.549 < 0.0000000000000002 ***
## regionQUEBEC   1.6486085   0.0943744  17.469 < 0.0000000000000002 ***
## regionATLANTIC  0.5126349   0.1818002   2.820    0.004806 **
## promo_units_per -7.0364858   0.5411102 -13.004 < 0.0000000000000002 ***
## altbev_units_per  7.2688615   0.9837629   7.389    0.0000000000000148 ***
## confect_units_per -0.8353451   1.1030602  -0.757    0.448871
## salty_units_per  26.0703631   1.4808940  17.604 < 0.0000000000000002 ***
## velocityA_units_per 11.4092559   3.1533882   3.618    0.000297 ***
## velocityB_units_per  5.1015780   3.2684432   1.561    0.118557
## velocityC_units_per 14.4423422   3.7934750   3.807    0.000141 ***
## velocityD_units_per  2.3609835   3.7725451   0.626    0.531424
## velocityNEW_units_per -33.5103228   5.0935739  -6.579    0.000000000047381 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 10471.8  on 7553  degrees of freedom
## Residual deviance:  7671.1  on 7540  degrees of freedom
## AIC: 7699.1
##
## Number of Fisher Scoring iterations: 5
```

Predicting the response variable on the test data using the training data

```
predict_test <- predict(model_train, newdata=data_test, type='response')
```

Forming table to look at the accuracy of the model

```
table2 <- table(predict_test>.5, data_test$high_med_units) #prediction on left and truth on top
my_confusion_matrix(table2)
```

```
##
##           0    1
##  FALSE 972 289
##   TRUE  314 942

## [[1]]
## [1] "942 = True Positive (TP), Hit"
##
## [[2]]
## [1] "972 = True Negative (TN), Rejection"
##
## [[3]]
## [1] "314 = False Positive (FP), Type 1 Error"
##
## [[4]]
## [1] "289 = False Negative (FN), Type 2 Error"
##
## [[5]]
## [1] "0.7604 = Accuracy (TP+TN/(TP+TN+FP+FN))"
##
## [[6]]
## [1] "0.7652 = Sensitivity, Recall, Hit Rate, True Positive Rate (How many positives did the model get right)"
##
## [[7]]
## [1] "0.7558 = Specificity, Selectivity, True Negative Rate (How many negatives did the model get right)"
##
## [[8]]
## [1] "0.7500 = Precision, Positive Predictive Value (How good are the model's positive predictions? TP/(TP+FP))"
##
## [[9]]
## [1] "0.7708 = Negative Predictive Value (How good are the model's negative predictions? TN/(TN+FN))"
```

Using the predictions above to help the business.

Putting the data back together for future use

```
# Put the prediction back into the test data
data_test$prediction <- predict_test

# Create a variable that shows if the prediction was correct
# (We have to do the classification--in `round(prediction)`--since logistic regression gives us a proba
data_test <- data_test %>%
  mutate(correct_prediction = if_else(round(prediction) == high_med_units, 'correct', 'WRONG!'))

# Add back the original data
temp1 <- ColumnsNotUsed[-partition, ]
full_test <- bind_cols(temp1, data_test)

# For viewing in class
full_test <- full_test %>%
  select(store, week, high_med_units, prediction, correct_prediction,
         size, region, promo_units_per, salty_units_per)
slice_sample(full_test, n=10)
```

```
## # A tibble: 10 x 9
##   store week high_med_units prediction correct_prediction size region
##   <dbl> <dbl>         <dbl>         <dbl> <chr>             <int> <fct>
## 1 63966   15             0         0.661 WRONG!             955 WEST
## 2 91361   23             1         0.692 correct            972 WEST
## 3 69921    3             0         0.405 correct            958 WEST
## 4 872     3             0         0.202 correct            937 ONTARIO
## 5 16110   21             0         0.528 WRONG!             915 QUEBEC
## 6 35155    9             0         0.0481 correct            930 ONTARIO
## 7 91226   14             1         0.893 correct            985 WEST
## 8 92576   49             1         0.910 correct           1015 WEST
## 9 37641    1             1         0.654 correct            967 WEST
## 10 77809   2             1         0.896 correct            982 WEST
## # i 2 more variables: promo_units_per <dbl>, salty_units_per <dbl>
```

##Summary of Findings

Feature/variable with the largest positive coefficient and statistically significant in the trained model summary

salty_units_per has the largest positive coefficient of 26.0703631 and is statistically significant because has a P-Value less than 0.05

Effect of selling a higher proportion of alternative beverages on the chance of having above median units sold

Selling a higher proportion of alternative beverages (variable “altbev_units_per”) **increases** the chance of having above median units sold. We know this because “altbev_units_per” has a coefficient of 7.2688615, which means that 1 unit increase in the “altbev_units_per” variable, will increase the median units sold by around 7.2688615.

Effect of selling a higher proportion of velocity B units on the chance of having above median units sold

Neither increase nor decrease the chance of having the above median units sold. We know this because the coefficient of “velocityB_units_per” is 5.1015780 but the associate p-value of 0.118557 is greater 0.05, which indicates that the coefficient for the said variable is not statistically significance. In addition, this

indicates that there is not enough evidence to conclude that selling a higher proportion of velocity B units will significantly affect, whether increase or decrease, the chance of having above median units sold.

Examining the accuracy of the predictions on the test data by answering whether there are more true positives or more true negatives.

There are more True Negatives (972) than True Positives (942).

First store in the ‘full_test’ dataset that has a “WRONG!” prediction

Store 186 located in ONTARIO region. This is true when stores are sorted by the ‘store’ feature in an ascending manner (lowest number first).

Why training data is used for the model training step

We use the **data-training** because it checks the quality and performance on how well the model generalizes new and unseen data.

Level of the variable is not present but accounted for in the intercept term

The level on the variable not present is but is accounted for in the intercept term is **ONTARIO**.

The feature ‘region’ has changed in the summary of the trained model. Further, only three regions show up in the summary of the model. The reasoning for this is that the ‘glm()’ function automatically recognizes that ‘region’ is a categorical variable (specifically a factor in R). Thus, the ‘glm()’ function has created “dummy variables” for the levels of ‘region’.

Interpreting the confusion matrix using the test/holdout data - highest value

The **Negative Predictive Value (NPV)** has the highest value of 0.7708. This means that the model is good at making predictions that a predicted negative outcome is actually negative. The model is 77.08% accurate when it comes to predicting negative outcome.

Interpreting the confusion matrix - lowest value

The **Precision, Positive Predictive Value** has the lowest value of 0.7500. This means that the model is not as good at making predictions that a predicted positive outcome is actually positive. The model is 75% accurate when it comes to predicting positive outcome.

Interpreting the confusion matrix - highest concern for NANSE

In NANSE’s business setting, the measure that NANSE care about the most is **Sensitivity**, which represents the proportion of weeks and stores with above-median sales that are correctly identified on the model. NANSE would care about **Sensitivity** because the store would want to ensure that the model will correctly capture instances of higher traffic and lessen the risk of false negatives. Meaning, NANSE will want to ensure that they minimize instances that their model fails to correctly predict above median sales when they occur, which could adversely impact business decisions in critical areas such as inventory management and resource allocation (e.g., marketing costs).